

# Project 2: Investigate TMDb movie data from kaggle

## Table of Contents

- [Introduction](#)
- [Data Wrangling](#)
- [Exploratory Data Analysis](#)
- [Conclusions](#)

## Introduction

The film industry is in full swing. The emergence of new multimedia platforms such as netflix or amazon prime, requires a redefinition and replacement of its structure and direction. The big Hollywood productions have a new competitor in a new format that day after day absorbs their expectations and customers. What are the characteristics of the sector? Is it really in decline? Throughout this document, we will prepare the database and describe the characteristics of this market. Using descriptive and inferential statistics we will try to establish a diagnosis of the sector

A study on the last years of film productions is an interesting exercise to understand the sector and the impact of new platforms on it. We propose to describe this market through certain problems, make a diagnosis and infer certain measures

To perform the analysis we will use the TMDb database This data set contains information about 10,000 movies collected from The Movie Database (TMDb), including user ratings and revenue.

- Certain columns, like 'cast' and 'genres', contain multiple values separated by pipe (|) characters.
- There are some odd characters in the 'cast' column. Don't worry about cleaning them. You can leave them as is.
- The final two columns ending with "\_adj" show the budget and revenue of the associated movie in terms of 2010 dollars, accounting for inflation over time.

## Problematics

- What are the main producers in the sector, in which genres did they invest the largest budgets? What results did they have?
- What are the most popular genres? How are they located in the valuation ranking? Are the most viewed the best valued too?
- On average, what is the most profitable genre? What is the least?
- What is the average budget for film productions? Is there a budget that ensures profitability?
- What are the movies with the most benefits, are they in the most popular genres?
- Is the scene dominated by a small number of directors or is there a diversity of production teams? Is there a director who ensures profitability?

In [1]:

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt

%matplotlib inline
```

## Data Wrangling

**DataBase reconnition:** In this section of the report, we will load in the data, check for cleanliness, and then trim and clean the dataset for analysis.

## General Properties


In [2]:

```
df = pd.read_csv(r'database/tmdb-movies.csv')
df_copy = df
df.head(3)
```

Out[2]:

	id	imdb_id	popularity	budget	revenue	original_title	cast
0	135397	tt0369610	32.985763	150000000	1513528810	Jurassic World	Chris Pratt Bryce Dallas Howard Irrfan Khan Vi...
1	76341	tt1392190	28.419936	150000000	378436354	Mad Max: Fury Road	Tom Hardy Charlize Theron Hugh Keays-Byrne Nic...
2	262500	tt2908446	13.112507	110000000	295238201	Insurgent	Shailene Woodley Theo James Kate Winslet Ansel...

3 rows × 21 columns



## Characteristics

to have a better understanding of the database characteristics, we perform some assessing descriptions of the data with the objective of build intuitions relatives to the **problematics** .

- First, we perform the functions shape and data types. to see the size of the dataset and the types of the columns. we notice that the column realise\_Date is in fact a date that need to be changed. with this information we can easily understand if the data is quantitative o categorical.
- Next, we perform the info function to relate the null values for each column and the data types
- With the function nunique, we see the number of differents values in each variable.
- Value counts is a really handy function to understand if a columns messy or not.

In [3]:

```
df.shape
```

Out[3]:

(10866, 21)

In [4]:

```
df.dtypes
```

Out[4]:

```
id                int64
imdb_id           object
popularity        float64
budget            int64
revenue           int64
original_title    object
cast              object
homepage          object
director          object
tagline           object
keywords          object
overview          object
runtime           int64
genres            object
production_companies object
release_date      object
vote_count        int64
vote_average      float64
release_year      int64
budget_adj        float64
revenue_adj       float64
dtype: object
```

In [5]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10866 entries, 0 to 10865
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                    10866 non-null  int64
1   imdb_id              10856 non-null  object
2   popularity           10866 non-null  float64
3   budget              10866 non-null  int64
4   revenue             10866 non-null  int64
5   original_title       10866 non-null  object
6   cast                10790 non-null  object
7   homepage            2936 non-null  object
8   director            10822 non-null  object
9   tagline             8042 non-null  object
10  keywords            9373 non-null  object
11  overview            10862 non-null  object
12  runtime             10866 non-null  int64
13  genres              10843 non-null  object
14  production_companies 9836 non-null  object
15  release_date        10866 non-null  object
16  vote_count          10866 non-null  int64
17  vote_average        10866 non-null  float64
18  release_year        10866 non-null  int64
19  budget_adj          10866 non-null  float64
20  revenue_adj         10866 non-null  float64
dtypes: float64(4), int64(6), object(11)
memory usage: 1.7+ MB
```

In [6]:

```
df.nunique()
```

Out[6]:

```
id                10865
imdb_id           10855
popularity        10814
budget            557
revenue           4702
original_title    10571
cast              10719
homepage          2896
director          5067
tagline           7997
keywords          8804
overview          10847
runtime           247
genres            2039
production_companies 7445
release_date      5909
vote_count        1289
vote_average       72
release_year       56
budget_adj         2614
revenue_adj        4840
dtype: int64
```

In [7]:

```
df.production_companies.value_counts().head(25)
```

Out[7]:

```
Paramount Pictures    156
Universal Pictures    133
Warner Bros.          84
Walt Disney Pictures  76
Columbia Pictures     72
Metro-Goldwyn-Mayer (MGM) 72
New Line Cinema       61
Touchstone Pictures   51
20th Century Fox     50
Twentieth Century Fox Film Corporation 49
TriStar Pictures      45
Orion Pictures        42
Miramax Films        32
Columbia Pictures Corporation 31
DreamWorks Animation  31
Pixar Animation Studios 30
Walt Disney Productions 29
Dimension Films       28
United Artists        23
Imagine Entertainment|Universal Pictures 22
Marvel Studios        22
The Asylum           21
Lions Gate Films      21
New World Pictures    17
Walt Disney Pictures|Pixar Animation Studios 17
Name: production_companies, dtype: int64
```

In [8]:

```
df.genres.value_counts().head(20)
```

Out[8]:

Drama	712
Comedy	712
Documentary	312
Drama Romance	289
Comedy Drama	280
Comedy Romance	268
Horror Thriller	259
Horror	253
Comedy Drama Romance	222
Drama Thriller	138
Comedy Family	102
Action Thriller	101
Thriller	93
Drama Comedy	92
Animation Family	90
Crime Drama Thriller	81
Crime Drama	74
Comedy Horror	72
Drama Comedy Romance	64
Action	63

Name: genres, dtype: int64

## ASSESS

- missing values: production companies, cast..
- inaccurate datatypes: release date
- inconsistency data: genres, production companies
- duplicated values
- the loc function is a perfect tool to filter the data and in this case look after some inconsistencies in it.

In [9]:

```
df.loc[df.duplicated(subset='original_title')].shape
```

Out[9]:

(295, 21)

In [10]:

```
df.loc[df.popularity != 0].shape
```

Out[10]:

(10866, 21)

In [11]:

```
df.loc[(df.revenue == 0) & (df.budget == 0)].shape
```

Out[11]:

(4701, 21)

## Data Cleaning

### Data type correction

In [12]:

```
df['release_date'] = pd.to_datetime(df.release_date)
```

### Dropping duplicated values

In [13]:

```
df.drop_duplicates(subset='original_title', keep='first', inplace=True)
```

### Reduce output values

In [14]:

```
df['genres'] = df['genres'].str.split('|', expand=True)[1]
```

In [15]:

```
df['production_companies'] = df['production_companies'].str.split('|', expand=True)[1]
```

### Inconsistency between values

In [16]:

```
def conditional_change(col_name, value_in, value_out):  
    df.loc[df[col_name].str.contains(value_in, case=False, na=False), col_name] = value_out  
    return df
```




In [17]:

```
conditional_change('production_companies', 'fox', 'Twentieth Century Fox Film Corporati  
on')  
conditional_change('production_companies', 'warner', 'Warner Bros')  
conditional_change('production_companies', 'disney', 'Walt Disney Production')
```

Out[17]:

	id	imdb_id	popularity	budget	revenue	original_title	cas
0	135397	tt0369610	32.985763	150000000	1513528810	Jurassic World	Chris Pratt Bryce Dallas Howard Irrfan Khan Vi.
1	76341	tt1392190	28.419936	150000000	378436354	Mad Max: Fury Road	Tom Hardy Charlize Theron Hugh Keays-Byrne Nic.
2	262500	tt2908446	13.112507	110000000	295238201	Insurgent	Shailen Woodley The James Kat Winslet Ansel.
3	140607	tt2488496	11.173104	200000000	2068178225	Star Wars: The Force Awakens	Harrison Ford Mark Hamill Carrie Fisher Adam D.
4	168259	tt2820852	9.335014	190000000	1506249360	Furious 7	Vin Diesel Paul Walker Jason Statham Michelle .
...	...	...	...	...	...	...	.
10861	21	tt0060371	0.080598	0	0	The Endless Summer	Michael Hynson Robert August Lord 'Tall Ho' B.
10862	20379	tt0060472	0.065543	0	0	Grand Prix	James Garner Eva Marie Saint Yves Montand Tosh.
10863	39768	tt0060161	0.065141	0	0	Beregis Avtomobilya	Innokenti Smoktunovskiy Oleg Efremov Georgi Z.
10864	21449	tt0061177	0.064317	0	0	What's Up, Tiger Lily?	Tatsuy Mihashi Akiko Wakabayashi Mi Hama Joh.
10865	22293	tt0060666	0.035919	19000	0	Manos: The Hands of Fate	Harold F. Warren Torrey Neyman John Reynolds Dian.

10571 rows × 21 columns



In [18]:

```
df.production_companies.value_counts().head(70)
```

Out[18]:

```
Warner Bros                224
Twentieth Century Fox Film Corporation    74
Touchstone Pictures        63
Metro-Goldwyn-Mayer (MGM)    62
Walt Disney Production     57
...
Robert Simonds Productions    9
British Broadcasting Corporation (BBC)    9
Lakeshore Entertainment      9
Stage 6 Films                9
Vertigo Entertainment         9
Name: production_companies, Length: 70, dtype: int64
```

### Innaccurates values (original title)


In [19]:

```
df.loc[df.duplicated(subset='original_title')]
```

Out[19]:

id	imdb_id	popularity	budget	revenue	original_title	cast	homepage	director	tagline
----	---------	------------	--------	---------	----------------	------	----------	----------	---------

0 rows × 21 columns



In [20]:

```
df.drop_duplicates(subset='original_title' , keep='first', inplace=True)
```

### Innaccurates values (release\_date)

there are 328 movies registered with the wrong year, to correct them we use an apply function combined with a lambda function to replace the values that are wrong

In [21]:

```
df.loc[df.release_date.dt.year > 2015].shape
```

Out[21]:

(328, 21)

In [22]:

```
df['release_date'] = df['release_date'].apply(lambda x: x.replace(year = x.year -100) if x.year>2015 else x)
```

the revenue and budget column have almost half of the registered values indicating 0. This represents almost half of the dataset. here we have different options:

- to use the mean of each year and each genre to approximate a value
- to consider them as null values and drop them.

In [23]:

```
df_2 = df.loc[(df.revenue != 0) & (df.budget != 0)]
```

- We generate a new dataframe df2 with complete observations for each row.
- This new df is generated to study the evolution and direction of the budget and revenues of the film industry over the last 60 years.

## Drop unnecessary columns

In [24]:

```
df_2.columns
```

Out[24]:

```
Index(['id', 'imdb_id', 'popularity', 'budget', 'revenue', 'original_title',
      'cast', 'homepage', 'director', 'tagline', 'keywords', 'overview',
      'runtime', 'genres', 'production_companies', 'release_date',
      'vote_count', 'vote_average', 'release_year', 'budget_adj',
      'revenue_adj'],
      dtype='object')
```

In [25]:

```
df_2.drop(columns=['imdb_id', 'cast', 'homepage', 'tagline', 'keywords', 'overview', 'runtime'], inplace=True)
```

c:\program files\python38\lib\site-packages\pandas\core\frame.py:3987: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
return super().drop()
```

In [26]:

```
df_2.head()
```

Out[26]:

	id	popularity	budget	revenue	original_title	director	genres	production
0	135397	32.985763	150000000	1513528810	Jurassic World	Colin Trevorrow	Adventure	Amblin
1	76341	28.419936	150000000	378436354	Mad Max: Fury Road	George Miller	Adventure	k
2	262500	13.112507	110000000	295238201	Insurgent	Robert Schwentke	Science Fiction	Ma
3	140607	11.173104	200000000	2068178225	Star Wars: The Force Awakens	J.J. Abrams	Adventure	Truenort
4	168259	9.335014	190000000	1506249360	Furious 7	James Wan	Crime	

Describe values

In [27]:

```
df_2.describe().T
```

Out[27]:

	count	mean	std	min	25%	50%
id	3756.0	4.055726e+04	6.786629e+04	5.000000	6.430000e+03	1.137050e+04
popularity	3756.0	1.194688e+00	1.486161e+00	0.001117	4.631333e-01	7.983430e-01
budget	3756.0	3.752862e+07	4.244601e+07	1.000000	1.000000e+07	2.400000e+07
revenue	3756.0	1.082386e+08	1.770478e+08	2.000000	1.358668e+07	4.483448e+07
vote_count	3756.0	5.307348e+02	8.857413e+02	10.000000	7.100000e+01	2.050000e+02
vote_average	3756.0	6.164004e+00	7.949363e-01	2.200000	5.675000e+00	6.200000e+00
release_year	3756.0	2.001538e+03	1.106193e+01	1960.000000	1.996000e+03	2.004000e+03
budget_adj	3756.0	4.436703e+07	4.498007e+07	0.969398	1.319180e+07	3.014881e+07
revenue_adj	3756.0	1.363931e+08	2.158969e+08	2.370705	1.830994e+07	6.092554e+07

## Exploratory Data Analysis

with the data wrangling done with use differents datasstes to find some insights. reated to the problematics that came out with the gather of the dataset.

**What was the evolution of the profits of the film industry during the period 1960 - 2018?**

to understand the evolution of the budget and the revenues of the industrie we use the function group by. and then, we plot it with a line plot using seaborn.

In [28]:


```
df_year = df_2.groupby(df['release_date'].dt.year).mean().reset_index()
```

In [29]:

```
def lineplot_x_2(x_axes, y1_axes, y2_axes, df):
    sns.set(style="darkgrid")
    plt.figure(figsize=(16,9))
    ax = sns.lineplot(x=x_axes, y=y1_axes, label=f"mean {y1_axes}", data=df)
    ax = sns.lineplot(x=x_axes, y=y2_axes, label=f"mean {y2_axes}", data=df)
    ax.set_title(f'Average {y1_axes} and {y2_axes} by year', fontsize=18)
    ax.set_xlabel('years')
    ax.set_ylabel('results');
```

In [30]:

```
lineplot_x_2('release_date', 'revenue', 'budget', df_year)
```



- the revenue curve is very intermittent. Growth processes last 2-3 years and are interrupted by abrupt falls. this pattern is reproduced throughout the study period
- at the end of the 80s, the amplitude of the volatility of the revenues is reduced and the difference between this curve and that of the budget that is maintained at around 450,000,000.00 is enlarged.
- The maximum recorded results are 1,600,000,000.00 relative to 2015, the last year with records.

In [31]:

```
df_year_sum = df_2.groupby(df['release_date'].dt.year).sum().reset_index()
df_year_sum['profit'] = df_year_sum['revenue'] - df_year_sum['budget']
df_year_sum.sort_values(by='profit', axis=0, ascending=False).head(10)
```

Out[31]:

	release_date	id	popularity	budget	revenue	vote_count	vote_average
55	2015	40164198	457.110856	7170777528	26202922801	155369	999.0
54	2014	32176813	416.406465	7393768000	24069969357	169922	1045.6
52	2012	12619538	222.580212	7587610525	24184456032	160629	946.3
51	2011	11301674	238.473152	8416550543	23404448100	121027	1203.1
53	2013	22144450	243.244332	8372922778	22974907074	177627	1097.4
49	2009	3425233	197.929015	7511827919	21563814641	102827	1046.3
50	2010	6327730	208.805990	8441108439	21802891959	115323	1074.1
48	2008	1824188	178.003082	6991556800	18599657891	85593	991.5
47	2007	1476451	159.612919	6740517892	18321062428	74054	980.8
44	2004	1267794	162.605879	6661999308	16381794281	71424	890.4

- In general terms both curves, are increasing. The revenue presents a higher volatility but the tendency shows how it is separating from the budget curve. In the last 20 years, there are two low bounces probably relates to the online piracy with web like piratebay or mega and in the earlier 2010 probably related with the new multimedia platforms.
- Even so, they revenues continue to grow in a faster than they budgets that stays leveled around the 400.000.000 for the all industrie
- To check for more granularity, we add the month value to the dataset, and we use the rolling function smooth the curves.

In [32]:

```
df_2['month_year'] = df_2['release_date'].dt.to_period('M')
df_y_m = df_2.groupby([df_2.month_year]).mean().reset_index()
df_y_m
```

<ipython-input-32-d9f477d823fc>:1: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)  
df\_2['month\_year'] = df\_2['release\_date'].dt.to\_period('M')

Out[32]:

	month_year	id	popularity	budget	revenue	vote_count	vote_
0	1960-06	284.000000	0.947307	3.000000e+06	2.500000e+07	235.000000	7
1	1960-10	966.000000	1.872132	2.000000e+06	4.905000e+06	224.000000	7
2	1960-12	18973.000000	0.055821	3.000000e+06	7.100000e+06	13.000000	7
3	1961-01	12230.000000	2.631987	4.000000e+06	2.158800e+08	913.000000	6
4	1961-03	18647.000000	0.173731	6.000000e+06	4.300000e+06	17.000000	6
...	...	...	...	...	...	...	...
531	2015-08	254943.916667	1.903474	3.291667e+07	7.925642e+07	659.416667	6
532	2015-09	271716.041667	1.804688	3.292083e+07	1.050920e+08	701.875000	6
533	2015-10	264129.529412	2.141399	3.543706e+07	9.103131e+07	710.764706	6
534	2015-11	266496.000000	2.278568	4.359444e+07	1.108521e+08	658.888889	6
535	2015-12	264217.333333	4.516088	7.333333e+07	3.860243e+08	1765.222222	6

536 rows × 10 columns



- before plotting the results, we use the rolling function to smooth the curve and make it more readable


In [33]:

```
df_y_m['budget_semestre'] = df_y_m['budget'].rolling(6).mean()
df_y_m['revenue_semestre'] = df_y_m['revenue'].rolling(6).mean()
```



In [34]:

```
lineplot_x_2(df_y_m.index, 'revenue_semestre', 'budget_semestre', df_y_m)
```



- The two curves presents some paralalism. The tendency shows that when a budget breaks the market a big increase in the revenue happens.the movements in the budget curve exponentially affect the profit curve
- The last 8 years, presents the budgets expends with intermittent evolution follow by an intermittent revenue that in generals terms is increassing. Lets' make a zoom of it using plt.xlimt

In [35]:

```
df_y_m.sort_values(by='revenue', axis=0,ascending=False)
```

Out[35]:


	month_year	id	popularity	budget	revenue	vote_count	vote_
136	1982-04	601.000000	2.900556	1.050000e+07	7.929106e+08	1830.000000	7
114	1980-01	1891.000000	5.488441	1.800000e+07	5.384000e+08	3954.000000	8
80	1975-06	578.000000	2.563191	7.000000e+06	4.706540e+08	1415.000000	7
535	2015-12	264217.333333	4.516088	7.333333e+07	3.860243e+08	1765.222222	6
463	2009-12	17482.090909	1.680397	7.281818e+07	3.816423e+08	1114.454545	6
...	...	...	...	...	...	...	...
91	1977-01	72277.000000	0.191541	8.100000e+05	1.258000e+06	12.000000	5
48	1970-02	2998.000000	0.279079	7.000000e+06	9.000000e+05	30.000000	6
44	1969-06	576.000000	0.615889	6.244087e+06	6.386410e+05	121.000000	7
134	1982-02	14373.000000	0.464188	2.000000e+00	1.600000e+01	27.000000	5
100	1978-04	28932.000000	0.439210	1.100000e+01	1.100000e+01	13.000000	6

536 rows × 12 columns



In [36]:

```
sns.set(style="darkgrid")
plt.figure(figsize=(16,9))
ax = sns.lineplot(x=df_y_m.index, y="revenue", label="mean revenue", data=df_y_m)
ax = sns.lineplot(x=df_y_m.index, y="budget", label="mean budget", data=df_y_m)
ax.set_title('Average revenue and budget by semestre', fontsize=18)
ax.set_xlabel('years')
ax.set_ylabel('results')
ax = plt.xlim(476, 535 )
```



en los ultmos 15 años, la curva del budget se mantiene con oscilaciones moderadas. que impacatan exponecialmente al revenue. en el 2009, se presenta un maximo, que vuelve a ser superado en 2012 y en 2015.

## Budget & revenues by genre

- After having seen the evolution of the values generated by the industry over time, we decided to identify the results of each of the genres, to have a more complete perspective of the viewer's tastes and the companies investments.
- The group by allow done it before, allows us to make this plot using seaborn box plot

In [37]:

```
def boxplot(x_axes, y_axes, df):  
    sns.set(style='darkgrid')  
    plt.figure(figsize=(16,9))  
    ax = sns.boxplot(x=x_axes, y=y_axes, orient= "h", data=df)  
    ax.set_title(f'{x_axes} by {y_axes}', fontsize=18)  
    ax.set_xlabel(f'{x_axes}')  
    ax.set_ylabel(f'{y_axes}');
```

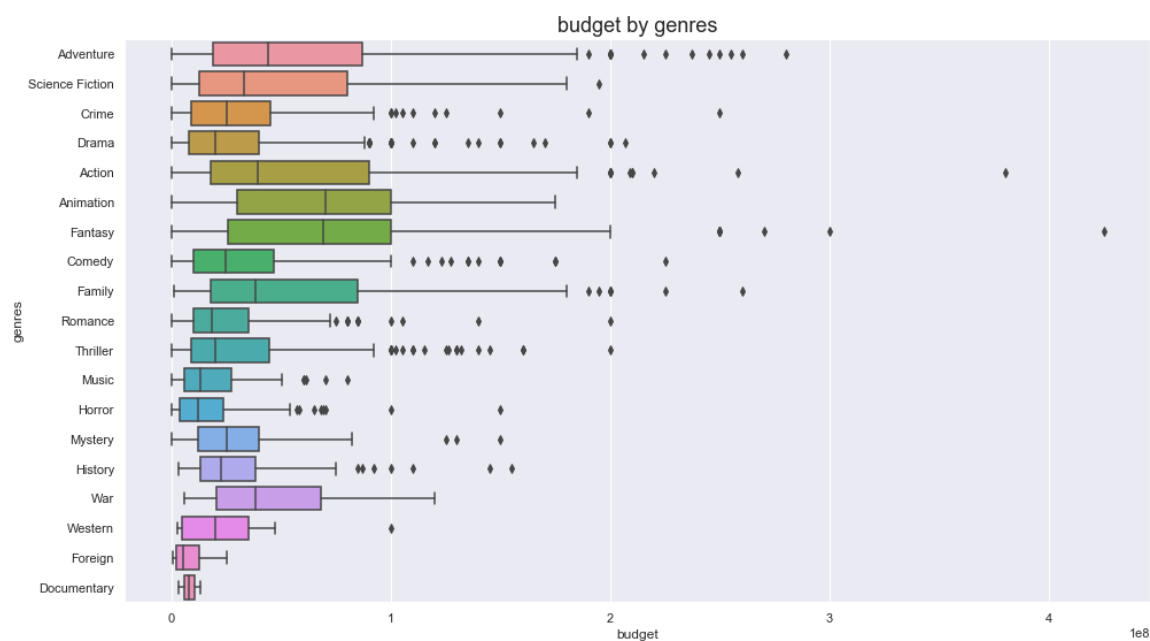
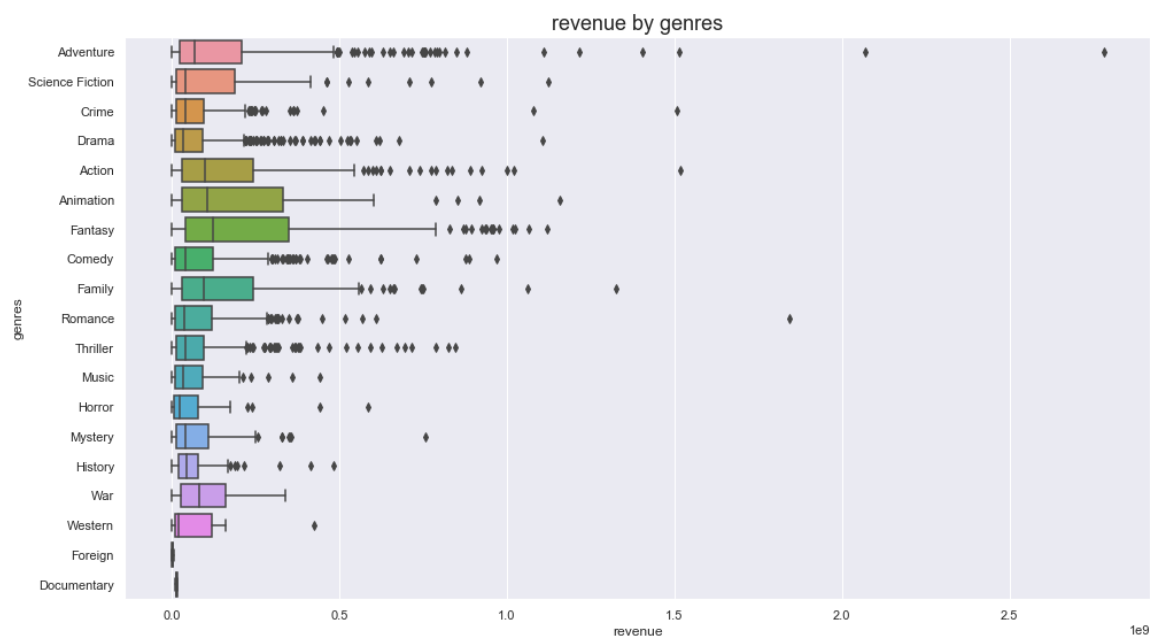
In [38]:

```

margins = ['revenue', 'budget']

for margin in margins:
    boxplot(margin, 'genres', df_2)

```



in average, the profits of the film industries are exponential.

the genres that requires more budgets to make a production are Fantasy and animation, follows by adventure, At the same time, its revenues are the highest the foreign productions and the documentary are bellow the average.

the diversity existing in the sector, has as a consequence a very large amplitude in monetary terms

## Budget & Revenue relationship

To understand the relationship between revenue and budget we can use the scatterplot.

In [39]:

```
def scatterplot(x_axes, y_axes, hue, df):
    sns.set(style='darkgrid')
    plt.figure(figsize=(16,9))

    ax = sns.scatterplot(x=x_axes, y=y_axes, hue=hue, data=df)
    ax.set_title(f'{x_axes} & {y_axes} by {hue}', fontsize=18)
    ax.set_xlabel(f'{x_axes}')
    ax.set_ylabel(f'{y_axes}');
```

In [40]:

```
scatterplot('budget', 'revenue', 'genres', df_2)
```



There is a moderate positive correlation between budget and revenue.

## Budget & Revenue by production companie

after having studied the behavior of the industry. we propose to focus on the main companies and observe their results. We know that he manages 22% of the profits, so his behaviors define the general results.

In [41]:

```
df_companies = df_2.groupby(by=['production_companies'], as_index=False).sum()
df_companies.sort_values(by='revenue', ascending=False).head(10)
```

Out[41]:

	production_companies	id	popularity	budget	revenue	vote_count	vo
1196	Warner Bros	4366930	200.362563	6819177000	18901681612	97912	
1151	Twentieth Century Fox Film Corporation	2042835	124.332429	3781350000	18724840150	71654	
862	Pixar Animation Studios	489436	56.476804	2301000000	9698819254	45111	
55	Amblin Entertainment	779493	84.065109	1834800000	9053944857	38989	
1168	Universal Pictures	880892	52.005811	2564777025	6587816806	20459	
780	New Line Cinema	238946	43.772675	1067000000	5988101890	29690	
821	Original Film	1152462	36.146177	1725000000	5962013012	29108	
1193	Walt Disney Production	1018681	49.080999	2303431000	5648410737	20170	
565	Jerry Bruckheimer Films	128277	31.591672	1855000000	5120058749	19287	
633	Legendary Pictures	665497	55.847220	1410000000	4216538498	30711	

In [42]:

```
profit_ten = df_companies.sort_values(by='revenue', ascending=False).head(10).iloc[:,4]
.sum() / df_2.iloc[:,3].sum()*100
```

```
print(f' the total revenue generated by the film industry ascends to {df_2.iloc[:,3].sum()}, of witch {profit_ten} % is controled by the 10 mayor companies in the industry ')
```

the total revenue generated by the film industry ascends to 406544127181, of witch 22.11376811378069 % is controled by the 10 mayor companies in the industry

In [43]:

```
prin_comp = list(df_companies.sort_values(by='revenue', ascending=False).head(10).iloc[:,0])
```

In [44]:

```
df_prin_comp = df_2.loc[df_2['production_companies'].isin(prin_comp)]
```

In [45]:

```
df_time_comp = df_prin_comp.groupby(['release_year', 'production_companies'], axis=0, as_index=False).sum()
df_time_comp.sort_values('revenue', ascending=False).head(10)
```

Out[45]:

	release_year	production_companies	id	popularity	budget	revenue	vote_c
135	2009	Twentieth Century Fox Film Corporation	114662	17.075648	603000000	4541406391	1
55	1997	Twentieth Century Fox Film Corporation	19468	7.697376	256500000	2242688658	
103	2004	Warner Bros	31789	11.683866	454000000	1682199952	
176	2015	Amblin Entertainment	431495	36.633973	190000000	1676139283	
178	2015	Original Film	425704	12.979555	248000000	1656420175	
114	2006	Twentieth Century Fox Film Corporation	5029	6.755179	290000000	1484709726	
169	2014	Legendary Pictures	331258	31.667504	395000000	1366357750	1
159	2012	Warner Bros	117760	8.594509	294500000	1313365415	
125	2008	Legendary Pictures	7995	10.308507	290000000	1267921825	
40	1993	Amblin Entertainment	753	4.582214	85000000	1241365768	

In [46]:

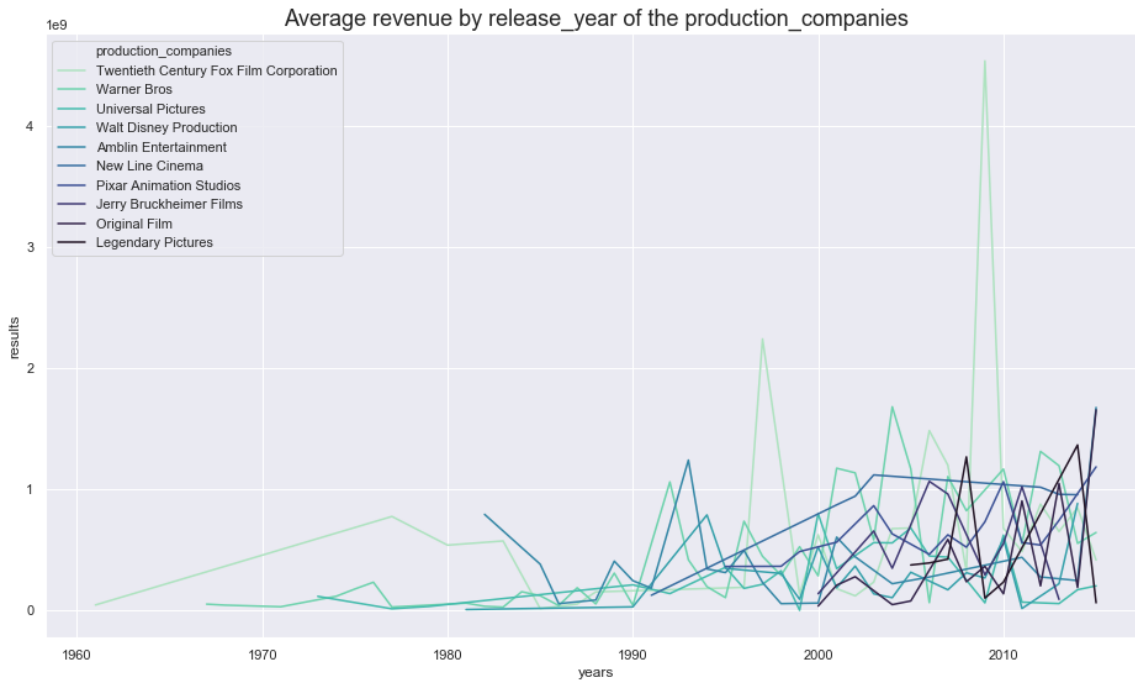
```
def lineplot_x_1(x_axes, y_axes, hue, df):
    sns.set(style="darkgrid")
    plt.figure(figsize=(16,9))
    palette = sns.color_palette("mako_r", df[f'{hue}'].nunique())
    ax = sns.lineplot(x=x_axes, y=y_axes, hue=hue, palette=palette, data=df)

    ax.set_title(f'Average {y_axes} by {x_axes} of the {hue}', fontsize=18)
    ax.set_xlabel('years')
    ax.set_ylabel('results');
```



In [47]:

```
lineplot_x_1('release_year', 'revenue', 'production_companies', df_time_comp)
```

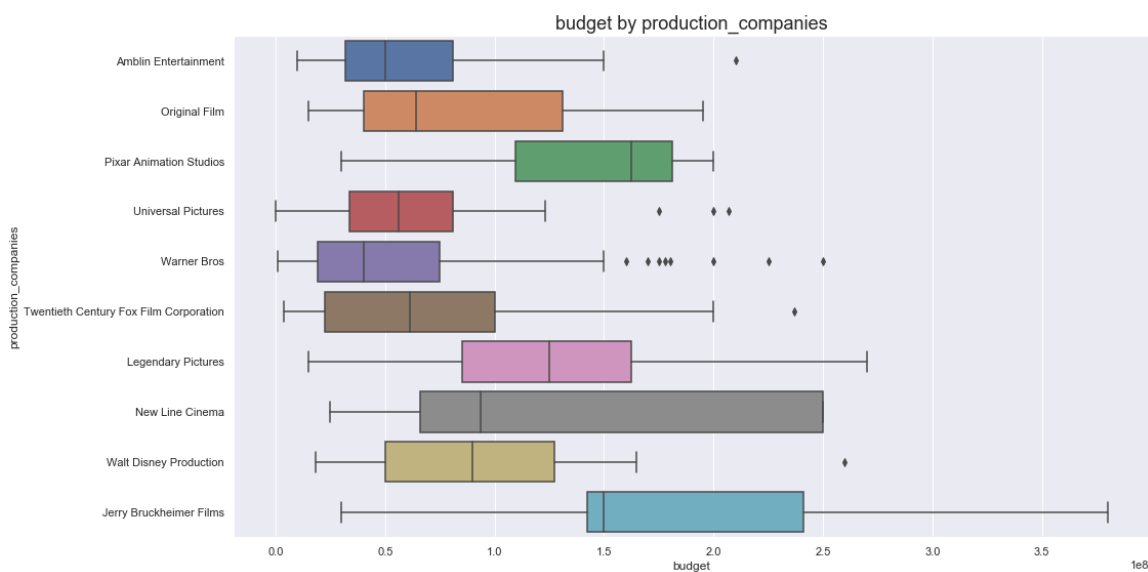
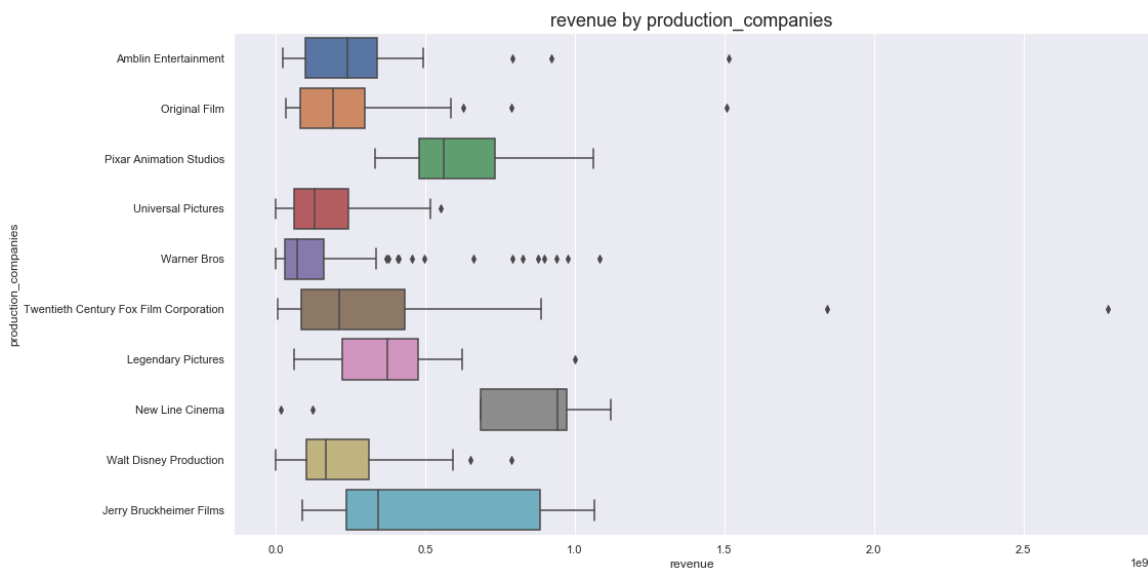


twenty century fox is the longest-running production company on the mayors 10 list. company profit curves show 2 extreme outlier. both belonging to twenty century fox and directed by the same director James Cameron: Titanic and avatar

Pixar animations is the cmost stable of the mayor ten.

In [48]:

```
for margin in margins:
    boxplot(margin, 'production_companies', df_prin_comp)
```



in average Pixar Animations is the company with most budgets investissements, clossely followed by Jerry Bruckheimer Films, but with a wider dstribution. in fact, this company spended the bigger budget for a film production on another hand, Warner Bros has the smaller average budget but with the biggest number of possitives outliers regarding the revenues, new line cinema leads the table. followed by Pixar animations, which has no adversaries in its genre the films that collected the most belong in first and second place to 20 century fox, and in third place to amban films

# What is the most popular genre?

the values related to the popularity and rating of the viewers was complete in the kaggle dataset, to answer this question then we use that dataset again as the largest number of values


In [49]:

```
df.head()
```

Out[49]:

	id	imdb_id	popularity	budget	revenue	original_title	cast
0	135397	tt0369610	32.985763	150000000	1513528810	Jurassic World	Chris Pratt Bryce Dallas Howard Irrfan Khan Vi...
1	76341	tt1392190	28.419936	150000000	378436354	Mad Max: Fury Road	Tom Hardy Charlize Theron Hugh Keays-Byrne Nic...
2	262500	tt2908446	13.112507	110000000	295238201	Insurgent	Shailene Woodley Theo James Kate Winslet Ansel... http:/
3	140607	tt2488496	11.173104	200000000	2068178225	Star Wars: The Force Awakens	Harrison Ford Mark Hamill Carrie Fisher Adam D...
4	168259	tt2820852	9.335014	190000000	1506249360	Furious 7	Vin Diesel Paul Walker Jason Statham Michelle ...

5 rows × 21 columns




In [50]:

```
#table = pd.pivot_table(df, values='D', index=['A', 'B'], columns=['C'], aggfunc=np.sum)
df.pivot_table(values='popularity', index=['genres'], columns=['release_year'], aggfunc=
=np.mean)
```

Out[50]:

release_year	1960	1961	1962	1963	1964	1965	1966	1967
genres								
Action	NaN	0.275009	1.391843	NaN	1.185667	1.910465	0.291704	0.438311
Adventure	1.872132	0.657446	0.468481	0.781120	NaN	0.312067	0.297781	0.285571
Animation	NaN	2.631987	NaN	NaN	NaN	NaN	0.276133	NaN
Comedy	NaN	NaN	0.153654	0.203298	0.544644	NaN	0.223819	0.101721
Crime	NaN	0.806519	NaN	NaN	0.626005	0.289649	0.737730	NaN
Documentary	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
Drama	0.412792	0.390316	0.453233	0.324263	0.231447	0.213986	0.263784	0.582211
Family	0.278064	0.280829	NaN	1.179561	0.665686	0.957326	0.540682	NaN
Fantasy	NaN	NaN	NaN	NaN	0.160098	NaN	NaN	NaN
Foreign	NaN	NaN	NaN	0.201754	0.152209	NaN	NaN	NaN
History	0.162753	0.538364	NaN	0.804533	0.321129	0.153438	0.418900	NaN
Horror	NaN	NaN	NaN	0.045417	0.349468	0.086219	0.034555	0.564781
Music	NaN	0.229252	0.323933	NaN	1.400006	0.624024	NaN	0.094541
Mystery	NaN	NaN	0.126140	0.408031	NaN	0.071354	0.509263	NaN
Romance	0.465497	0.200269	0.124418	0.331643	0.162341	0.209034	0.239435	0.275151
Science Fiction	NaN	0.194495	0.179212	0.467836	0.237514	0.199141	0.408574	0.242831
TV Movie	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
Thriller	0.446302	NaN	0.308033	1.113477	0.371976	0.213382	0.402730	1.554801
War	NaN	NaN	NaN	NaN	NaN	0.246816	0.317824	NaN
Western	0.268273	NaN	0.209355	NaN	NaN	0.141221	0.267152	NaN

20 rows × 56 columns




In [51]:

```
df_popularity = df.groupby(['release_year', 'genres'], as_index=False).mean()
df_popularity['pop_semestre'] = df_popularity['popularity'].rolling(12).mean()
```

In [52]:

```
lineplot_x_1('release_year', 'pop_semestre', 'genres', df_popularity)
```



Against the first impressions, the audience popularity increase exponentially. The last 5 years of the record registered have been the most attractive to general audiences

In [61]:

```
df_popularity.loc[df_popularity['release_year']==2015].sort_values(by='popularity',ascending=False)
```

Out[61]:

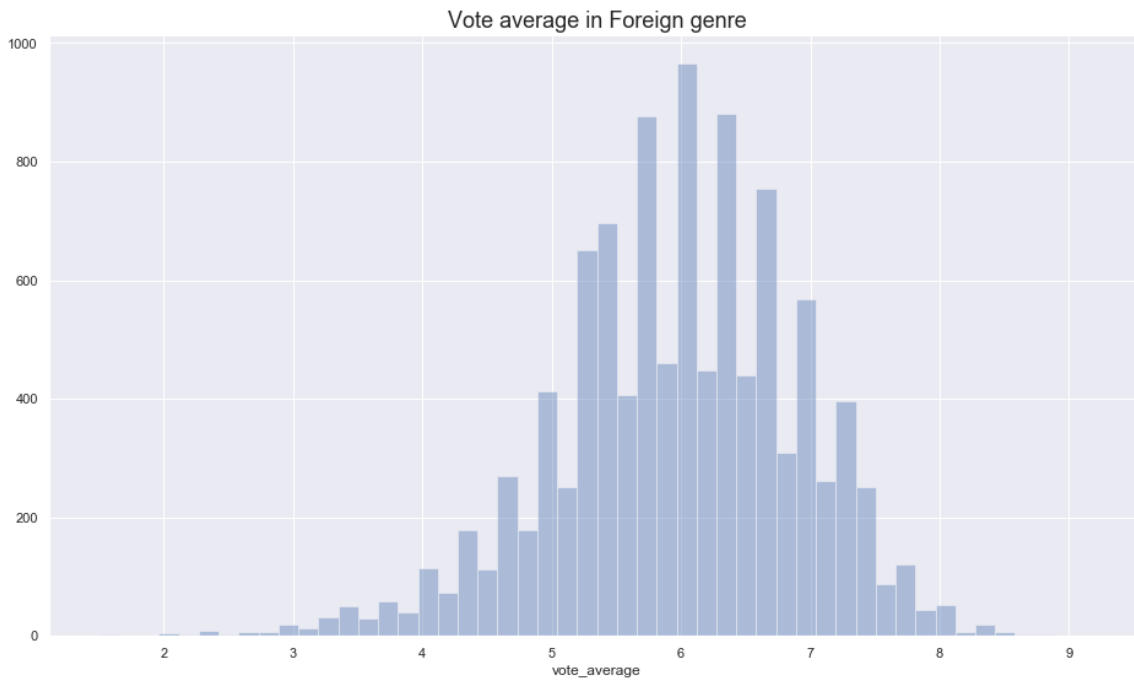
	release_year	genres	id	popularity	budget	revenue	rt
854	2015	Adventure	251646.200000	4.665070	6.216000e+07	3.140688e+08	101.1
861	2015	Fantasy	209785.500000	3.533745	7.075000e+07	1.815848e+08	116.5
857	2015	Crime	287416.125000	1.906528	3.525000e+07	1.218250e+08	116.6
855	2015	Animation	274308.687500	1.903577	4.131250e+07	2.065514e+08	80.2
860	2015	Family	289652.800000	1.808794	6.800000e+07	6.748480e+07	84.6
867	2015	Science Fiction	289027.346154	1.344678	1.025000e+07	2.626538e+07	92.8
853	2015	Action	297687.289474	1.315581	1.685921e+07	4.569690e+07	92.8
859	2015	Drama	294297.709302	1.189173	1.160814e+07	3.243085e+07	104.9
866	2015	Romance	279791.785714	1.065802	4.714286e+06	4.568167e+07	106.6
871	2015	Western	321466.000000	0.993217	3.090000e+07	0.000000e+00	125.5
865	2015	Mystery	294790.461538	0.904213	5.346154e+06	1.483834e+07	100.6
869	2015	Thriller	287017.693878	0.865156	5.869388e+06	1.641517e+07	94.8
856	2015	Comedy	290402.210526	0.842056	1.226276e+07	3.974254e+07	88.0
864	2015	Music	302290.944444	0.739470	5.522222e+06	2.869487e+07	107.1
863	2015	Horror	309882.590909	0.529663	2.834091e+06	8.733089e+06	89.1
862	2015	History	289493.400000	0.491164	6.260000e+06	3.200484e+06	103.6
858	2015	Documentary	327465.909091	0.279312	0.000000e+00	1.729417e+06	90.8
870	2015	War	325258.000000	0.238515	3.000000e+06	4.538850e+04	83.5
868	2015	TV Movie	358140.545455	0.221470	2.727273e+05	0.000000e+00	96.0



In [70]:

```
#genres = [i for i in df.genres.unique().tolist() if type(i) == str]
#for genre in genres:

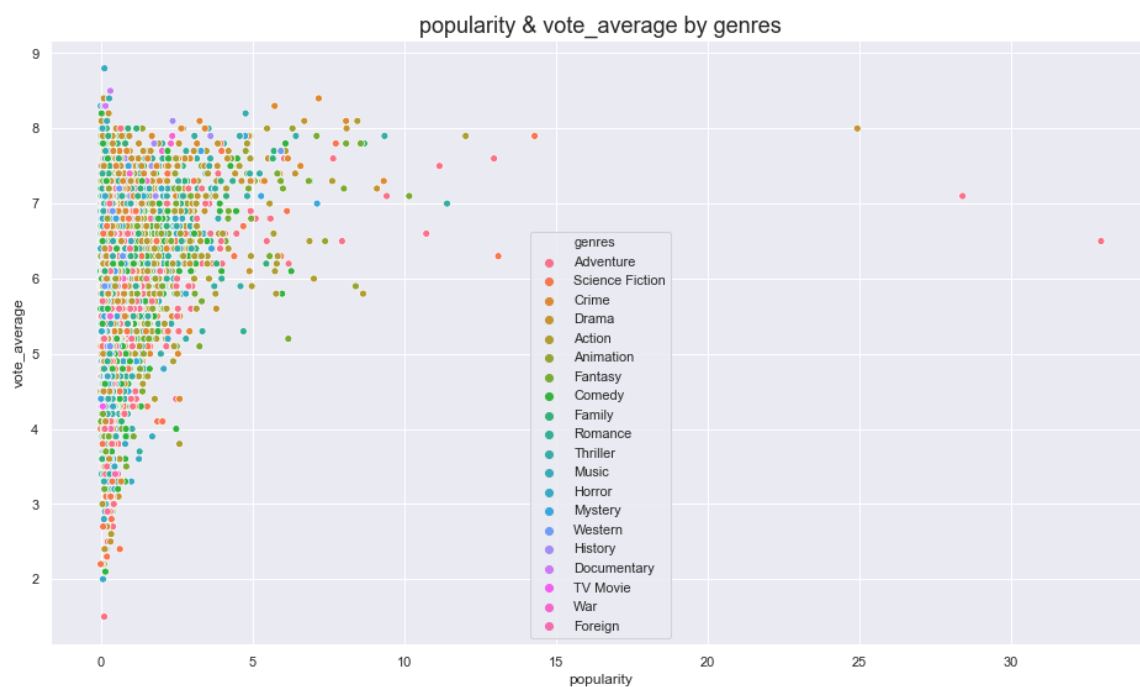
sns.set(style='darkgrid')
plt.figure(figsize=(16,9))
plot=df['vote_average']
ax = sns.distplot(plot,kde=False,hist=True)
ax.set_title(f'Vote average in {genre} genre', fontsize=18);
```



less than 1000 movies have been rates with more than 7 points  
over 1500 have been rated with 5 or less. over 4500 have been rated between 5.5 and 6.5

In [55]:

```
scatterplot('popularity', 'vote_average', 'genres', df)
```



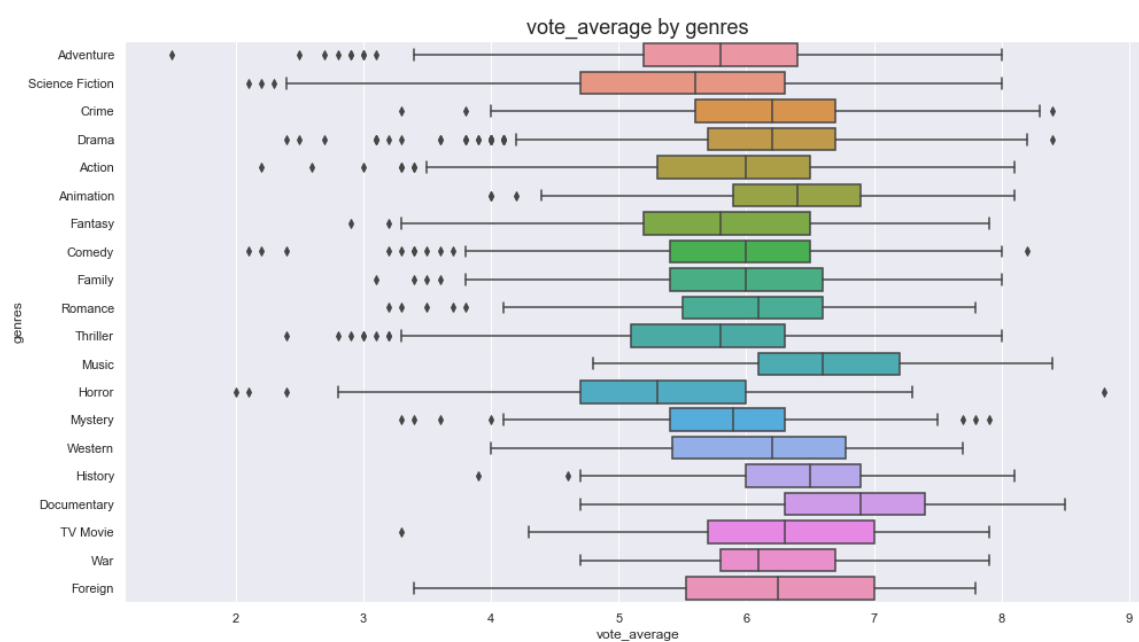
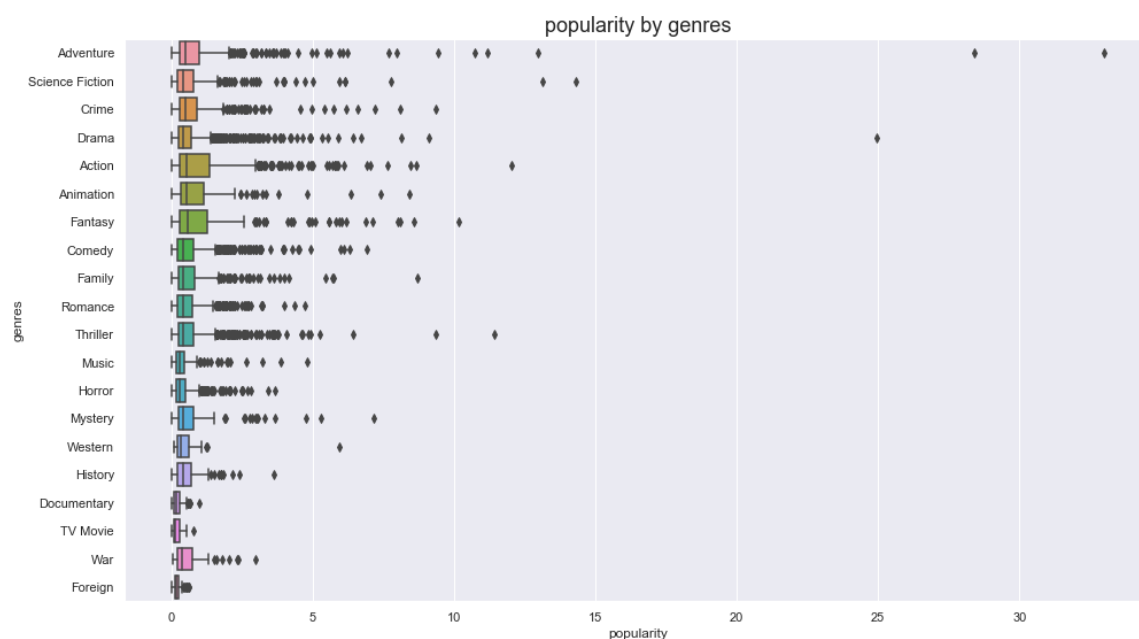
Dispite the blockbuster outliers. A strong positive correlation exists between the popularity and the average votes of a movie



In [56]:

```
people_choice = ['popularity', 'vote_average']

for choice in people_choice:
    boxplot(choice, 'genres', df)
```



the genre most popular is action, but the most populars movies are in the adventure genre.  
 the amount of outliers affects the visualization of this variable  
 the votes average is the result of what it seems an online survey. the volumen of votes registers is significantly small in comparison with the audience popularity. so this results can be bias.  
 The big majority of genres has the average votes between 6 and 7.

## Conclusions

- despite the new competitors that have emerged in recent years and the increase in illegal online download and playback platforms. The profits of the great companies in the film industry are exceeded year by year, with the virtue of having leveled expenses on average.
- the volatility of the benefits is directly related to the budgets invested in that period. movements in the budget curve have an exponential impact on the profit curve
- the total revenue generated by the film industry ascends to 406.544.127.181,00 us, of witch 22.11 % is controled by the 10 mayor companies.
- The adventure genre presents two big hits but in general terms the audiences preferes action or fantasy. Also comedy over drama
- in valoration termes, the betters are music film or documentary boths requiered a slow budget but they don't appeal to the big audiences. this cases can presents a bias in the votes

### limitations:

Related to the monetary values of the film industry we find certain limitations.

- Budgets and revenues do not have currency units, we must assume that they are in it to perform the analysis.
- The time factor in the budgets was not taken into account, inflation coefficients from different periods were not considered. So the values at the temporal level can have partial results.
- Within the monetary variables we find both budgets and benefits with values equal to zero (almost half of the dataset). To do the study, we discarded them from the sample and considerably reduced the number of elements.

Regarding the ratings and ratings of the films

- Absence of spatial information. The lack of information in this regard prevents us from clustering or generating a diagnosis about any spatial unit. They are absolute values without spatial reference.
- Very uneven relationship between popularity and number of votes. The opinion of the voters is much lower in scale than the number of spectators, the values are at risk of being partial

In [57]:

```
from subprocess import call  
call(['python', '-m', 'nbconvert', 'Investigate_a_Dataset.ipynb'])
```

Out[57]:

4294967295